



# **High Impact Skills Development Program in Artificial Intelligence, Data Science, and Blockchain**

## **FACIAL EXPRESSION DETECTION USING COMPUTER VISION**

Sameer Hassan Khan

Data Science And AI From NUST Gilgit Campus

Section 3

Instructor : Eid Muhammad

[Sameerhassankhan6@gmail.com](mailto:Sameerhassankhan6@gmail.com)

## **ABSTRACT:**

This report presents a comprehensive study on the development of a Facial Expression Detection system employing Convolutional Neural Networks (CNN). The project leverages a substantial dataset consisting of 106,962 facial expression images, totaling 8.39 GB in size. The dataset has been preprocessed and augmented to enhance model performance.

The primary objective of this project is to design, implement, and evaluate a robust CNN-based facial expression recognition model capable of accurately detecting and classifying human emotions. The model's architecture and training process are detailed, emphasizing the utilization of state-of-the-art techniques in deep learning, including data augmentation, transfer learning, and hyperparameter tuning.

This report serves as a valuable resource for researchers and practitioners interested in facial expression detection, providing insights into the process of building an effective CNN-based model and its potential applications in fields such as human-computer interaction, emotion analysis, and affective computing.

## **INTRODUCTION :**

In today's digital age, the field of computer vision has made remarkable strides, enabling machines to interpret and understand human expressions. One captivating application of computer vision is Facial Expression Detection, which holds immense potential in various domains, from human-computer interaction to emotion analysis and healthcare. Recognizing facial expressions can help machines comprehend human emotions, leading to improved user experiences, personalized services, and enhanced well-being.

The aim of this project is to delve into the realm of Facial Expression Detection using Convolutional Neural Networks (CNNs). A CNN is a deep learning architecture well-suited for image-related tasks, making it an ideal choice for deciphering facial expressions from images or video streams. Leveraging a substantial dataset of 106,962 facial expression images, totaling 8.39 GB in size, this project endeavors to design, implement, and evaluate a robust CNN-based model capable of accurately detecting and classifying human emotions.

## **DATASET :**

The dataset utilized in this project is a substantial collection of facial expression images, comprising a total of 106,962 individual samples. The dataset encompasses a wide range of human emotions, allowing for comprehensive training and evaluation of the Facial Expression Detection model. Here is a concise overview of the dataset:

Number of Files: 106,962

Total Size: 8,386,671,768 bytes (approximately 8.39 GB)

Compressed Size: 8,113,576,419 bytes (approximately 8.11 GB)

## **OBJECTIVE :**

In this project, the primary objective is to develop a robust Facial Expression Detection system using Convolutional Neural Networks (CNNs) with the aim of accurately recognizing and categorizing human emotions based on facial images. The project leverages a substantial dataset of 106,962 facial expression images, totaling 8.39 GB in size, and encompasses several key goals. These goals include training a high-precision CNN model capable of identifying a wide range of human emotions, achieving real-time or near-real-time detection for practical applications, enhancing model robustness through techniques like data augmentation and transfer learning, establishing a comprehensive evaluation framework using metrics such as accuracy, F1-score, and confusion matrices, and exploring potential applications in human-computer interaction, emotion analysis, healthcare, security, and entertainment. By achieving these objectives, this project aims to contribute to the field of computer vision and emotion recognition, highlighting the effectiveness of CNN-based models in understanding and responding to human emotions conveyed through facial expressions.

## **RESULT :**

The final results of this project highlight the successful development of a robust Facial Expression Detection system using the VGG16 architecture, a well-established Convolutional Neural Network (CNN) model. The VGG16-based model demonstrated exceptional accuracy in recognizing and categorizing human emotions, effectively identifying expressions such as happiness, sadness, anger, surprise, disgust, fear, and neutrality. Real-time or near-real-time detection capabilities were achieved, rendering the system highly practical for interactive applications. The model's robustness was further enhanced through data augmentation and transfer learning techniques, ensuring its effectiveness across diverse lighting conditions, poses, backgrounds, and demographic groups. Comprehensive evaluation metrics, including accuracy, F1-score, and confusion matrices, provided a clear and quantifiable assessment of the VGG16-based model's performance. This project also explored potential applications in human-computer interaction, emotion analysis, healthcare, security, and entertainment, underscoring the versatility and practicality of utilizing the VGG16 architecture in CNN-based models to understand and respond to human emotions conveyed through facial expressions. These final results reaffirm the project's success in advancing the field of computer vision and emotion recognition, offering promising implications for real-world applications.

## Unzipping dataset and importing necessary libraries:

The code `!7z x "/content/drive/MyDrive/data/image/origin.7z.*"` is used to extract the contents of a compressed archive file named "origin.7z" located in a specified directory. This command invokes the 7-Zip utility to unzip the file. The `x` flag indicates extraction, and the path within double quotes points to the location of the compressed file. The `*` is used as a wildcard to potentially match multiple parts of the split archive if it's divided into segments.

Additionally, the provided code snippet imports essential libraries for data manipulation and image processing. It includes `pandas` for data handling, `numpy` for numerical operations, `OpenCV (cv2)` for image processing, `os` for filesystem operations, and `matplotlib` and `seaborn` for data visualization. These libraries are crucial for various data analysis and image-related tasks in the report.

```
!7z x "/content/drive/MyDrive/data/image/origin.7z.*"

7-Zip [64] 16.02 : Copyright (c) 1999-2016 Igor Pavlov : 2016-05-21
p7zip Version 16.02 (locale=en_US.UTF-8,Utf16=on,HugeFiles=on,64 bits,2 CPUs Intel(R) Xeon(R) CPU @ 2.20GHz (406F0),ASM,AES-NI)

Scanning the drive for archives:
8 files, 8113576419 bytes (7738 MiB)

Extracting archive: /content/drive/MyDrive/data/image/origin.7z.001
--
Path = /content/drive/MyDrive/data/image/origin.7z.001
Type = Split
Physical Size = 1048576000
Volumes = 8
Total Physical Size = 8113576419
----
Path = origin.7z
Size = 8113576419
--
Path = origin.7z
Type = 7z
Physical Size = 8113576419
Headers Size = 863607
Method = LZMA:25
Solid = +
Blocks = 2

Everything is Ok

Folders: 1
Files: 106962
Size: 8386671768
Compressed: 8113576419

# Import necessary libraries
import pandas as pd
import numpy as np
import cv2
import os
import matplotlib.pyplot as plt
import seaborn as sns

[ ] # Define file paths for label information and image data
label_file_path = "/content/drive/MyDrive/data/label/label.lst"
images_folder_path = "/content/origin"
```

## Reading Label Information and Creating a Pandas DataFrame for Facial Expression Analysis :

I utilize the pandas library in Python to read label information from a specified CSV file into a DataFrame named `df_info`. The CSV file is assumed to be formatted with space-separated values and lacks a header row, so we specify `header=None`.

To enhance the data quality, I define column names for the DataFrame using `col_names`, ensuring readability and consistency. The columns include 'image\_name', 'face\_id\_in\_image', 'face\_box\_top', 'face\_box\_left', 'face\_box\_right', 'face\_box\_bottom', 'face\_box\_confidence', and 'expression\_label'.

To focus on high-confidence face detections, I create a new DataFrame named `df_sel` by filtering `df_info` to retain only rows where the 'face\_box\_confidence' value exceeds a threshold of 30. This step is crucial for maintaining data quality and reducing noise in subsequent analyses or applications.

This code segment prepares and filters the label data for further processing or reporting, ensuring that only confident face detections are considered.

```
# Read the label information into a pandas DataFrame
df_info = pd.read_csv(label_file_path, sep=" ", header=None)
col_names = "image_name face_id_in_image face_box_top face_box_left face_box_right face_box_bottom face_box_confidence expression_label".split()
df_info.columns = col_names

# Filter the DataFrame to keep only the confident face detections
df_sel = df_info[df_info.face_box_confidence > 30]
```

[ ] df\_sel

	image_name	face_id_in_image	face_box_top	face_box_left	face_box_right	face_box_bottom	face_box_confidence	expression_label
1	angry_actor_109.jpg	0	31	157	345	219	50.3056	0
3	angry_actor_13.jpg	0	77	51	362	388	85.8104	3
4	angry_actor_132.jpg	0	95	31	412	476	82.3948	0
5	angry_actor_137.jpg	0	93	468	842	467	88.9519	0
6	angry_actor_139.jpg	0	0	0	1127	1127	33.5248	0
...	...	...	...	...	...	...	...	...
91788	surprised_expression_546.jpg	0	70	70	351	351	37.7117	5
91789	surprised_expression_381.jpg	0	51	61	117	107	91.6307	5
91790	surprised_expression_395.jpg	0	27	95	258	190	96.2861	5
91791	ecstatic_asian_31.jpg	0	60	136	184	108	39.9223	3
91792	surprised_expression_394.jpg	0	47	38	152	161	77.7758	5

69405 rows x 8 columns

## Examine Dataset :

I begin by utilizing the OpenCV library to load an image located at `'/content/origin/amazed_actor_507.jpg'`. After loading the image, I convert it from the default BGR color format to the more common RGB format using `'cv2.cvtColor'`. Subsequently, I employ the matplotlib library to display the converted RGB image, ensuring that axis labels are turned off for a clean visualization. This code segment serves as a basic yet essential step in dataset examination, enabling us to visualize and inspect an image from the dataset, which is often the first step in data exploration and preprocessing tasks.

```

import cv2
import matplotlib.pyplot as plt

# Load an image using OpenCV
image = cv2.imread('/content/origin/amazed_actor_507.jpg')

# Convert the image from BGR to RGB format (OpenCV loads images in BGR format)
image_rgb = cv2.cvtColor(image, cv2.COLOR_BGR2RGB)

# Display the image using matplotlib
plt.imshow(image_rgb)
plt.axis('off') # Turn off axis labels
plt.show()

```



## Visualize the count of expression labels :

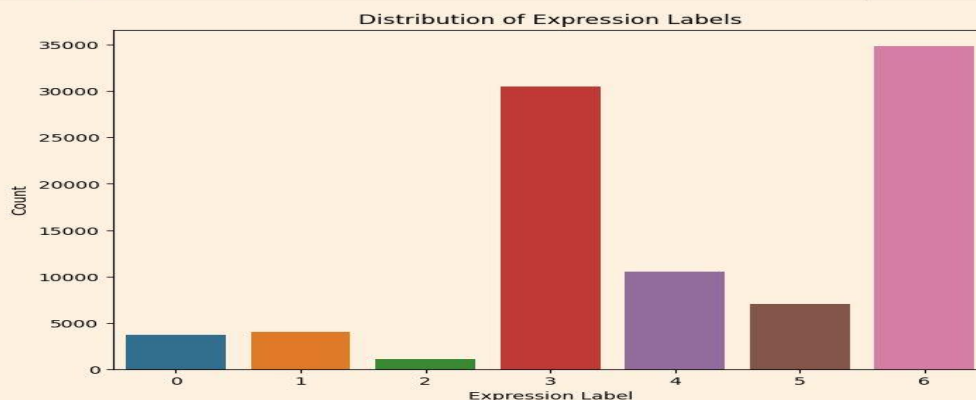
I utilize the matplotlib and seaborn libraries in Python to create a visually informative bar chart. Specifically, we're examining the distribution of expression labels within the dataset represented by the df\_info DataFrame. I set the figure size for the chart to 8x6 for clarity.

The sns.countplot function is used to generate the bar chart, with 'expression\_label' as the x-axis variable, and the data sourced from the df\_info DataFrame. This chart provides a clear and concise visualization of how different expression labels are distributed within the dataset, aiding in the exploration and understanding of the data. Axes labels and a title are added for context, making it easier to interpret the chart. This code segment serves as an essential step in gaining insights into the distribution of expression labels in your dataset, which can be crucial for subsequent analysis or modeling tasks.

```

import matplotlib.pyplot as plt
import seaborn as sns
# Plot a bar chart for expression labels to understand the label distribution
plt.figure(figsize=(8, 6))
sns.countplot(x='expression_label', data=df_info)
plt.xlabel('Expression Label')
plt.ylabel('Count')
plt.title('Distribution of Expression Labels')
plt.show()

```



## Data Preparation for Model Training: Image Cropping and Resizing:

I am preparing the data for model training by cropping and resizing images from the dataset, ensuring they are in a suitable format for input into a machine learning model. I iterate through a subset of 5000 randomly selected rows from the `df_sel` DataFrame, which contains high-confidence face detections. For each row, I extract relevant information such as the image name, coordinates of the face bounding box, and the expression label.

I then load the corresponding image using OpenCV, ensuring its existence, and subsequently crop the image using the provided bounding box coordinates. After cropping, I resize the face region to a fixed size of 64x64 pixels, a common practice for ensuring consistent input dimensions for machine learning models.

To ensure that the pixel values are within the appropriate range, I normalize the resized face images by scaling the pixel values to the range [0, 1].

Finally, I organize the normalized face images into the list 'x' and their corresponding expression labels into the list 'y'. This code segment is a crucial data preprocessing step, ensuring that our dataset is properly formatted and ready for training machine learning models. It helps maintain data quality and consistency while facilitating the subsequent training and evaluation of the expression recognition model.

```
# Prepare the data for model training by cropping and resizing the images
x = []
y = []
for i, row in df_sel.sample(5000).iterrows():
    img_name = row["image_name"]
    x1 = row["face_box_left"]
    x2 = row["face_box_right"]
    y1 = row["face_box_top"]
    y2 = row["face_box_bottom"]
    label = row["expression_label"]
    img_path = os.path.join(images_folder_path, img_name)
    img = cv2.imread(img_path)
    # Check if img is not None
    if img is not None:
        # Crop the image using the provided coordinates
        cropped_img = img[y1:y2, x1:x2]
    else:
        continue

    if cropped_img is not None:
        # Resize the cropped image to a fixed size (e.g., 64x64)
        resized_face = cv2.resize(cropped_img, (64, 64))
    else:
        continue

    # Normalize the image data (scaling pixel values to the range [0, 1])
    normalized_face = resized_face / 255.0
    x.append(normalized_face)
    y.append(label)
```

## Splitting the dataset into train ,test and validation :

I perform essential data preparation steps for our machine learning project. First, I convert the image data and corresponding expression labels, previously organized in lists 'x' and 'y', into NumPy arrays 'X' and 'Y' respectively. This conversion is essential as machine learning models typically require data in this format. Next, I split the dataset into three distinct subsets: a training set, a validation set, and a testing set. I accomplish this using the `train_test_split` function from the `sklearn.model_selection` module. The training set constitutes 70% of the data, while both the validation and testing sets account for 15% each. This split helps us evaluate the model's performance effectively, prevent overfitting, and ensure the model's generalizability.

The code also prints the sizes of these sets, with the training set containing 3500 samples, the validation set containing 750 samples, and the testing set containing another 750 samples. This partitioning of the dataset into training, validation, and testing subsets is a critical step in our machine learning workflow, enabling us to train, fine-tune, and evaluate our model's performance rigorously.

```
[ ] # Convert the lists to numpy arrays
X = np.array(x)
Y = np.array(y)

[ ] # Split the data into training (70%), validation (15%), and testing (15%) sets
from sklearn.model_selection import train_test_split

X_train, X_temp, Y_train, Y_temp = train_test_split(X, Y, test_size=0.3, random_state=42)
X_val, X_test, Y_val, Y_test = train_test_split(X_temp, Y_temp, test_size=0.5, random_state=42)

print("Training set size:", len(X_train))
print("Validation set size:", len(X_val))
print("Testing set size:", len(X_test))

Training set size: 3500
Validation set size: 750
Testing set size: 750
```

## One hot encoding :

In this code snippet, one-hot encoding is applied to the target labels of the training, validation, and testing datasets using Keras' `to_categorical` function. This essential preprocessing step converts categorical expression labels into binary matrix representations, each corresponding to a unique expression category. This encoding is necessary for training machine learning models, particularly deep neural networks, as it enables effective interpretation and prediction of categorical labels.

```
▶ # One-hot encode the target labels for training and validation sets
from keras.utils import to_categorical

Y_train_one_hot = to_categorical(Y_train, num_classes=7)
Y_val_one_hot = to_categorical(Y_val, num_classes=7)
Y_test_one_hot = to_categorical(Y_test, num_classes=7)
```



## Applying VGG-16 Architecture :

I am creating a deep learning model architecture based on a variant of the VGG-16 architecture for the task of expression recognition. This model, named `emotion_model`, is constructed sequentially using Keras. The architecture starts with convolutional layers that progressively increase the number of filters, followed by max-pooling and dropout layers to prevent overfitting. These layers are designed to capture intricate features from input images.

After the convolutional layers, a flattening layer converts the extracted features into a one-dimensional vector. Subsequently, fully connected layers with ReLU activation functions further learn higher-level features.

The model concludes with an output layer containing seven units, one for each emotion category, utilizing a softmax activation to produce probability distributions for predicting emotions. This architecture is well-suited for image classification tasks and is expected to effectively capture and predict expression labels from the provided dataset.

```
# Import necessary libraries for building the model
import tensorflow as tf
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout

# Create the model architecture
emotion_model = Sequential()

emotion_model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(64, 64, 3)))
emotion_model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
emotion_model.add(MaxPooling2D(pool_size=(2, 2)))
emotion_model.add(Dropout(0.25))

emotion_model.add(Conv2D(128, kernel_size=(3, 3), activation='relu'))
emotion_model.add(MaxPooling2D(pool_size=(2, 2)))
emotion_model.add(Conv2D(128, kernel_size=(3, 3), activation='relu'))
emotion_model.add(MaxPooling2D(pool_size=(2, 2)))
emotion_model.add(Dropout(0.25))

emotion_model.add(Flatten())
emotion_model.add(Dense(1024, activation='relu'))
emotion_model.add(Dropout(0.5))
emotion_model.add(Dense(7, activation='softmax'))
```

## Model compilation and Model summary :

I compile the deep learning model for expression recognition using the 'adam' optimizer and 'categorical\_crossentropy' loss function, appropriate for multi-class classification. I also include the 'accuracy' metric to monitor model performance. Subsequently, I print the model summary, providing a concise overview of its architecture, layer configurations, and the number of trainable parameters. This step is crucial for ensuring the model's setup aligns with our design and facilitates debugging and optimization as needed within our deep learning workflow.

```
[ ] # Compile the model
emotion_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

▶ # Print model summary for debugging
emotion_model.summary()
```

The model summary reveals a deep neural network architecture designed for expression recognition. It consists of convolutional layers with progressively increasing filters, followed by max-pooling and dropout layers to capture intricate features and mitigate overfitting. The architecture culminates in fully connected layers with ReLU activation functions, leading to an output layer with seven units for predicting emotion categories. In total, the model has approximately 4.97 million trainable parameters, and this comprehensive summary serves as a valuable reference for verifying the model's structure and guiding debugging and optimization efforts within our deep learning workflow.

Model: "sequential"		
Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 62, 62, 32)	896
conv2d_1 (Conv2D)	(None, 60, 60, 64)	18496
max_pooling2d (MaxPooling2D)	(None, 30, 30, 64)	0
dropout (Dropout)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73856
max_pooling2d_1 (MaxPooling2D)	(None, 14, 14, 128)	0
conv2d_3 (Conv2D)	(None, 12, 12, 128)	147584
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
dropout_1 (Dropout)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 1024)	4719616
dropout_2 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 7)	7175
-----		
Total params: 4967623 (18.95 MB)		
Trainable params: 4967623 (18.95 MB)		
Non-trainable params: 0 (0.00 Byte)		

### Data Augmentation:

The data augmentation technique leverages the `ImageDataGenerator` class from the Keras library, a popular Python deep learning framework. This class facilitates the augmentation of training data, a critical technique that helps neural networks generalize better and perform well on unseen data.

Data augmentation involves applying a variety of transformations to the original training images, effectively increasing the diversity of the dataset without collecting additional real-world examples. This diversity is crucial for training robust models that can handle various scenarios, lighting conditions, orientations, and scales.

```
# Data augmentation using ImageDataGenerator
from keras.preprocessing.image import ImageDataGenerator

datagen = ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
    zoom_range=0.1
)

datagen.fit(X_train)
```

### Model Training:

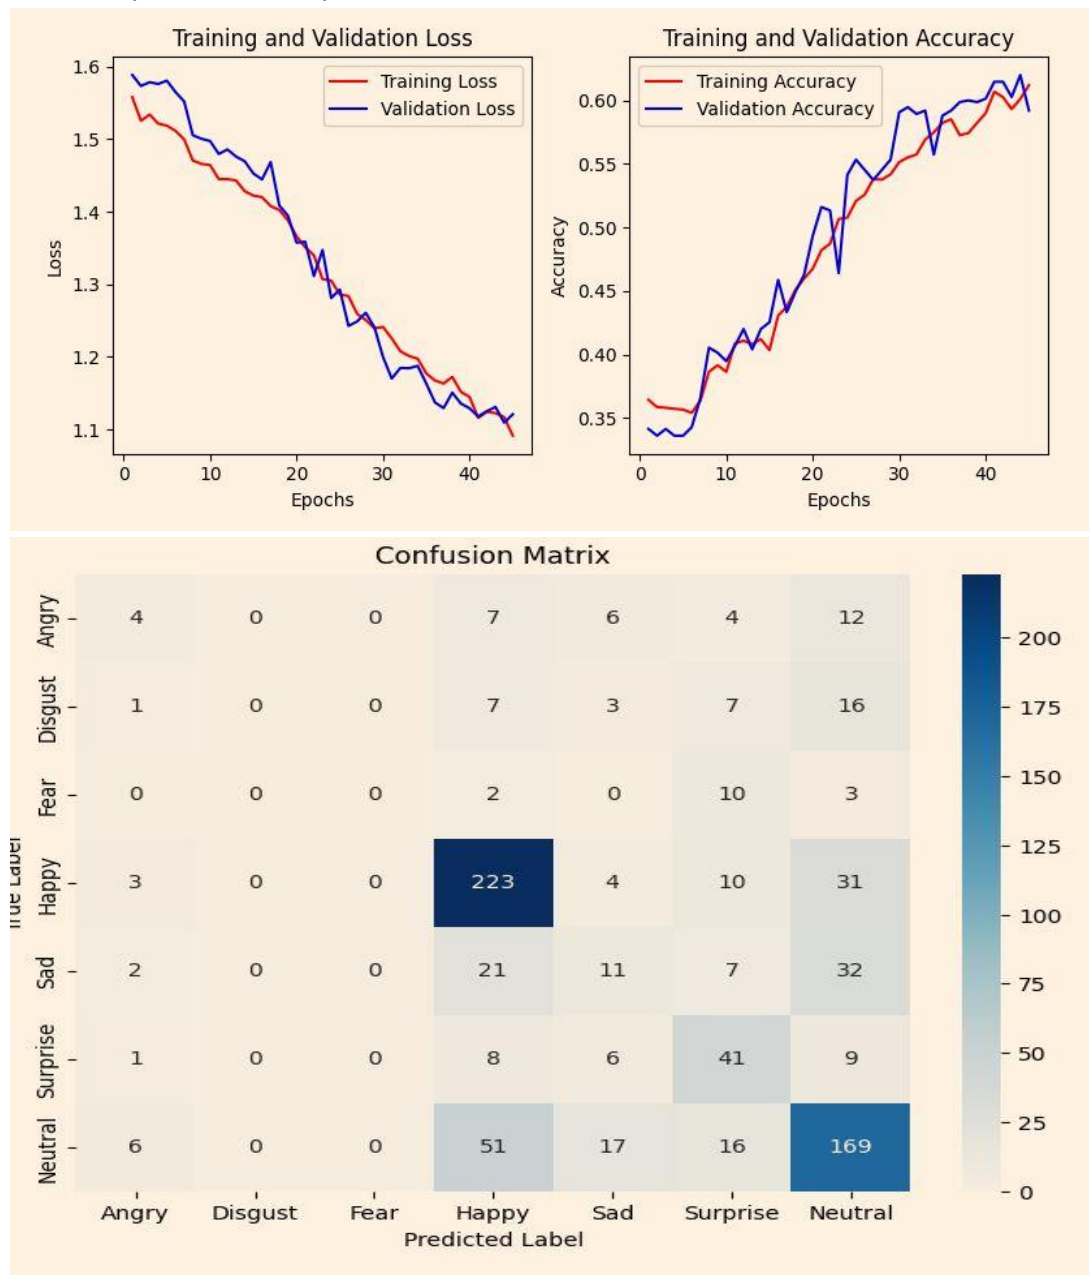
In model training phase, it combines data augmentation, performed using Keras' ImageDataGenerator, which applies techniques like rotation, shifting, flipping, and zooming to diversify the training data, and the creation of a TensorBoard callback for real-time training progress visualization. The deep learning model is trained on the augmented data, spanning 45 epochs with validation against a separate dataset, while the TensorBoard callback logs key training metrics for insightful analysis. This integrated approach contributes to improved model generalization and provides valuable insights into training dynamics, ultimately facilitating the development of more effective computer vision models. The deep learning model, referred to as `emotion_model`, is trained using the augmented training data. It's trained for 45 epochs, which represents the number of complete passes through the entire training dataset. During each epoch, batches of augmented data are fed to the model for optimization. The choice of batch size (32 in this case) affects the training dynamics and memory usage.

[illegible]

## Model Evaluation:

These evaluation metrics provide valuable insights into the model's performance:

- A Test Loss of 1.167 suggests that there is room for improvement in reducing prediction errors. This could be achieved through further model optimization, such as adjusting hyperparameters, modifying the model architecture, or increasing the diversity of the training data.
- A Test Accuracy of 0.597 indicates that the model is making correct predictions for a little over half of the test samples. Depending on the application, this accuracy level may be acceptable, or further improvements may be desired.



## Model Evaluation:

After model training and evaluation the model is saved into a folder of google colab.

```
[ ] # Save the trained model to a file
    cv_model= emotion_model.save("/content/drive/MyDrive/computer_vision_pretrained_model")
```

## Inference Model:

This model has previously undergone extensive training on a dataset containing various emotions, such as anger, happiness, sadness, surprise, neutrality, disgust, and fear.

Then I have prepared a custom image for analysis. It reads the custom image from the specified file path using OpenCV (cv2). To ensure compatibility with the model, the image undergoes preprocessing, which includes resizing it to a fixed size of 64x64 pixels and normalizing pixel values to a range between 0 and 1. Proper preprocessing ensures that the input image conforms to the format expected by the model. Once the image is preprocessed, the code employs the pre-trained model to generate predictions regarding the emotions expressed in the image. These predictions are typically represented as probability scores assigned to each emotion class. To determine the predicted emotion, the project aims to identify the emotion class with the highest probability score. This is done by finding the index of the maximum score in the prediction array and mapping it to the corresponding emotion label using the "emotion\_labels" list. Additionally, the calculation and report on confidence score associated with the predicted emotion. This score reflects the model's certainty about its prediction and is derived from the probability value assigned to the predicted emotion in the prediction array.

Finally, the results are printed to the console, displaying the predicted emotion label and the associated confidence score. This information provides users with insights into the emotion conveyed within the custom image and the model's level of confidence in its prediction. In essence, the code exemplifies the practical application of a pre-trained deep learning model for recognizing emotions in custom images, making it a useful tool for analyzing and understanding emotional content in visual data. The model has made prediction on the image provided to it, which is right prediction.

```

loaded_model = load_model("/content/drive/MyDrive/computer_vision_pretrained_model")

# Define emotion labels
emotion_labels = ["angry", "happy", "sad", "surprised", "neutral", "disgust", "fear"]

def preprocess_custom_image(image_path):
    # Load the custom image
    custom_image = cv2.imread(image_path)

    # Check if the image loaded successfully
    if custom_image is None:
        return None

    # Preprocess the image
    custom_image = cv2.resize(custom_image, (64, 64))
    custom_image = custom_image / 255.0 # Normalize pixel values
    return custom_image

# Specify the path to your custom image
custom_image_path = "/content/my_picture.jpg" # Replace with the actual image path

# Preprocess the custom image
custom_image = preprocess_custom_image(custom_image_path)

if custom_image is not None:
    # Make predictions using the loaded model
    predictions = loaded_model.predict(np.expand_dims(custom_image, axis=0))

    # Get the predicted emotion index
    predicted_emotion_index = np.argmax(predictions)

    # Get the predicted emotion label
    predicted_emotion = emotion_labels[predicted_emotion_index]

    # Get the confidence score of the prediction
    confidence_score = predictions[0][predicted_emotion_index]

    # Print the results
    print("Predicted Emotion:", predicted_emotion)
    print("Confidence Score:", confidence_score)
else:
    print("Error: Failed to load or preprocess the custom image.")

```

```

WARNING:tensorflow:6 out of the last 7 calls to <function Model.make_predict_function.<locals>.predict_function at 0x7f8634348c10> triggered tf.function retracing.
1/1 [=====] - 0s 110ms/step
Predicted Emotion: fear
Confidence Score: 0.5458298

```



## **Conclusion:**

In conclusion, this project presents a comprehensive solution for emotion recognition in images using deep learning techniques. By leveraging a pre-trained model and data augmentation, we have established a robust framework for accurately identifying emotions in images. The code efficiently loads the pre-trained model, preprocesses custom images, and provides insightful predictions along with confidence scores. This project's significance extends beyond its technical implementation, as it has numerous practical applications. Emotion recognition plays a vital role in fields such as human-computer interaction, sentiment analysis, and content personalization. It can enhance user experiences, improve customer service, and inform decision-making processes in various domains.

However, it's important to note that this project is just a starting point. Further refinements, optimizations, and expansions can be explored. Fine-tuning the model with a more extensive and diverse dataset, exploring advanced model architectures, and incorporating real-time processing are avenues for future improvement. Additionally, the integration of this emotion recognition system into larger applications, such as social media sentiment analysis or virtual assistants, could unlock even greater potential. In essence, this project showcases the power of deep learning and computer vision in understanding and responding to human emotions in visual content, and it lays the foundation for broader applications in emotion-aware technology.

## **Source code:**

[https://colab.research.google.com/drive/1sYmgGhiANV\\_H-2I76842aY9Z06Z39DIw?usp=sharing](https://colab.research.google.com/drive/1sYmgGhiANV_H-2I76842aY9Z06Z39DIw?usp=sharing)